




UNIVERSITY OF  
CAMBRIDGE

# Grapt: A Hybrid Database System for Flexible Retrieval of Graph-structured Data

Christopher J. O. Little



Emmanuel College

*A dissertation submitted to the University of Cambridge  
in partial fulfilment of the requirements for the degree of  
Master of Engineering in Computer Science*

University of Cambridge  
Computer Laboratory  
William Gates Building  
15 JJ Thomson Avenue  
Cambridge CB3 0FD  
UNITED KINGDOM

Email: [cl554@cam.ac.uk](mailto:cl554@cam.ac.uk)

May 30, 2016



# Declaration

I, Christopher J. O. Little of Emmanuel College, being a candidate for Part III of the Computer Science Tripos, hereby declare that this dissertation and the work described in it are my own work, unaided except as may be specified below, and that the dissertation does not contain material that has already been used to any substantial extent for a comparable purpose.

Total word count: 11001

**Signed:**

**Date:**

This dissertation is copyright ©2016 Christopher J. O. Little.

All trademarks used in this dissertation are hereby acknowledged.



# Abstract

There has recently been an increasing need for fast analysis of graph-structured data, which has led to the development of several graph-centric alternatives to traditional relational databases. Although these ensure the fast execution of queries which fit within this graph-centric model, inevitably a compromise has been made, and other queries perform less efficiently than they would have done with a relational database. I propose the development of a query language enabling intelligent dispatch of optimised queries either directly to a relational database, or through a graph-centric query processing layer.



# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Research Goals . . . . .	2
<b>2</b>	<b>Background</b>	<b>5</b>
2.1	Relational Databases . . . . .	5
2.2	Graph Databases . . . . .	6
2.3	Related Work . . . . .	7
<b>3</b>	<b>Implementation</b>	<b>11</b>
3.1	System Overview . . . . .	11
3.2	Row Handler . . . . .	13
3.3	Graph Handler . . . . .	13
3.3.1	Lookahead Prefetcher . . . . .	14
3.3.2	Block Prefetcher . . . . .	15
3.3.3	Iterative Lookahead Prefetcher . . . . .	15
3.4	Graph API . . . . .	17
3.5	Query Processor . . . . .	21
3.5.1	Improving on the Graph Manipulation Kernel . . . . .	23
3.5.2	Executing gSQL . . . . .	25
<b>4</b>	<b>Evaluation</b>	<b>27</b>
4.1	Test Framework . . . . .	28
4.2	Prefetcher Performance . . . . .	29
4.3	Query Performance . . . . .	40
4.3.1	Performance of Graph-centric Queries . . . . .	41
4.3.2	Performance of Row-Centric Queries . . . . .	42
4.3.3	Performance of Hybrid Queries . . . . .	45
4.3.4	Expressiveness of Query Systems . . . . .	47
<b>5</b>	<b>Conclusion</b>	<b>51</b>





# List of Figures

4.1	Average Execution time for A* search of different path lengths using a CTE lookahead prefetcher . . . . .	30
4.2	Average Execution time for A* search of different path lengths using an iterative lookahead prefetcher . . . . .	31
4.3	Average Execution time for A* search of different path lengths using a block prefetcher . . . . .	32
4.4	Graph miss rate for A* search of different path lengths using a CTE lookahead prefetcher . . . . .	33
4.5	Graph miss rate for A* search of different path lengths using an iterative lookahead prefetcher . . . . .	34
4.6	Graph miss rate for A* search of different path lengths using a block prefetcher . . . . .	35
4.7	Average cost of cache miss for A* search of different path lengths using a CTE lookahead prefetcher . . . . .	36
4.8	Average cost of cache miss for A* search of different path lengths using an iterative lookahead prefetcher . . . . .	37
4.9	Average cost of cache miss for A* search of different path lengths using a block prefetcher . . . . .	38
4.10	Average A* Execution time for paths of varying length across three database engines . . . . .	43
4.11	Average Levenshtein-search execution time for 100 queries across three database engines . . . . .	44
4.12	Average Execution time for 100 queries consisting of a Min. Levenshtein-distance search and an A* path search across three database engines. . . . .	46



# List of Tables



# Chapter 1

## Introduction

Graph databases are an alternative to traditional Relational Database Systems (RDBMS), which improve access speeds for some queries by prioritising access to related entities, rather than expecting traversal through a fixed schema. This design can more naturally represent relationships between data in many fields, from the analysis of gene networks in biology to social network recommendation engines. It can provide huge performance improvements for graph-centric queries, such as shortest path queries.

For other queries, however – such as aggregation of otherwise unrelated data items – it is far more efficient to rely on a known schema over a fixed set of rows for fast access. There are many trade-offs to be made between the RDMS and Graph Database model, but the core division is that where an RDBMS is optimised for aggregated data, a graph database is optimised for highly connected data. In practice, however, although it is certainly the case that some datasets are more highly connected than others, many of today's business operations cannot be classified according to that metric alone.

Social networks are often used as an important example of a highly connected dataset today. However the reality is that, although relationships between friends are ideally modelled as a graph, this is not the case across the board. Performing queries across all users – such as discovering users who live within

particular geographic bounds – is much less efficient when every single edge between users needs to be considered during a graph traversal. This heterogeneity, coupled with the common uncertainty at design-time about which part of an application will be most frequently used, makes it difficult to choose an ideal database system for any particular application.

## 1.1 Research Goals

The optimisations made to increase the performance of graph databases for certain problem necessarily sacrifice performance for others. Thus it seems optimistic to aim to create a system which performs both as well as Neo4J for graph-centric problems and as well as PostgreSQL for row-centric problems. Instead, the aims for this project were to improve on the worst-case performance of these systems, while still minimising the loss in best-case performance. By doing this, we can provide a system which provides a higher average-case performance in the situation where graph-centric queries and row-centric queries are both used. More concretely, there were four goals for this project. Two are fairly trivial improvements over existing systems, while the other two arise precisely because of the hybrid nature of the system. Note that for this research project, we limit our attention to data-retrieval queries only.

### **1: Increase performance of graph queries compared to a relational engine**

Relational database engines perform poorly for graph-centric queries such as path-finding. A hybrid engine should perform better than this, even if it cannot achieve performance on par with purely graph-centric engines.

## **2: Increase performance of row queries compared to a graph engine**

Conversely, graph engines perform poorly for row-centric queries such as aggregation over all nodes, and a hybrid engine should aim to perform better than this.

## **3: Increase performance of hybrid queries compared to both engines**

Some queries may have multiple components to them, some of which would best be handled by a graph engine, while others best by a relational system. For example, we may wish to identify two users who are geographically closest to one another, and find a path between them through the social graph. Finding geographical neighbours involves comparisons between all nodes, not just ones which are local to one another in the social graph. This would thus be well handled by a row-centric system, and less well by a graph-centric one. The path-finding component, however, would perform much better under the graph system. Under either system, some component of the query would necessarily perform badly and cause a performance bottleneck. A hybrid system ought to not be affected so badly by either bottleneck, and may thus outperform both systems for this kind of query.

## **4: Expose a coherent interface for the hybrid system**

Relational databases have a well-established query interface in the form of *SQL*. Queries in *SQL* aim to create and select particular rows from the database, and are thus inherently row-centric, making it difficult for a programmer to even express graph-centric queries. Although the graph database community have not yet settled on a single query language to act as *SQL*'s graph-centric counterpart, the two most popular languages – Gremlin and

Cypher – suffer from the opposite problem: they are not well suited to expressing row-centric queries considering all nodes in the graph. Since a hybrid system must aim to be equally well suited in terms of performance to either sort of query, it should not give precedence to either in terms of interface. Graph queries should be as easy to express using the hybrid system as row-centric ones.

TODO: Mention that I’m only interested here in read-query ability, and that updates or insertions are not covered (In future work we can say that these shouldn’t be too difficult to add some sort of cache coherency model)

In the next chapter, I will first give a brief introduction to the technologies involved in both relational and graph databases, including their relative strengths and weaknesses. I will continue by providing a brief overview of the surrounding research landscape, particularly focussing on previous attempts to bring together graph and row-based systems. Chapter 3 discusses the implementation of a research prototype aiming to satisfy all four of my research goals: *GraphT*. The success of this prototype is experimentally evaluated in Chapter 4, along with direct comparisons between Neo4J and PostgreSQL. These comparisons serve not only to provide a point of reference to measure GraphT against, but also to validate the claims made above about the relative performance of the two systems in satisfying different types of query. Finally, Chapter 5 summarises the findings made, and discusses possible directions for future research.



# Chapter 2

## Background

### 2.1 Relational Databases

The relational database model first arose in in the 1970s with Codd’s [1] famous “Relational Model of Data.” The success of this model was largely due to its simplicity, and it quickly found favour with the industry at the time. Since then, the model has been carefully refined, with numerous implementations each offering incremental optimisations.

Along with simplicity, most RDBMS provide desirable guarantees such as atomicity, consistency, isolation and durability (ACID). These properties were formally defined by Gray [3], and are now highly prized since they give a high degree of confidence in the integrity of data, and allow databases to be parallelised and distributed across several machines without risk of conflict between multiple concurrent queries.

These factors have lead to widespread adoption of the relational database model, and a corresponding proliferation of implementations, including MySQL, SQL Server and PostgreSQL. I will be evaluating PostgreSQL in this report, and using it as a base relational store upon which to build a graph-centric hybrid store. PostgreSQL is a popular, open-source database system supporting most of the SQL standard. It also supports a procedural language

called PL/pgSQL which extends SQL to offer more programming control. In particular, PostgreSQL supports recursive SQL queries which are valuable for efficiently finding paths through an edge relation without performing exhaustive JOINS.

## 2.2 Graph Databases

Although Graph databases have enjoyed a recent resurgence in popularity, the principles themselves are not much younger than Codd’s relational database model. Indeed, both models aim to represent relationships between entities. Chen [2] provided one of the earliest incarnations in the form of his “Entity-relationship Model.” This found popularity alongside the development of object-oriented programming in the 1980s, and several variations to this model were presented at the time [7][4]. Kunii [5] presented the first properly-called graph database model, called G-Base. Interest waned in the late 90s for a variety of reasons, including a general move by the database research community towards semistructured data. Recent times have seen a resurgence in popularity, as alternative strategies are explored to handle the “Big Data” explosion in data size and complexity.

The need for graph representations of data can be seen through the number of large technology companies who have recently developed in-house solutions to the problem. In 2010, Facebook released Open Graph, and Twitter released FlockDB[14], both projects aiming to make traversal through the social networks more efficient. In 2012 Google revealed their Knowledge Graph [9], aiming to connect search terms to real-world entities. The benefit of these graph databases is that queries can be localised to a particular area within the graph. Without direct links between related entities, every entity in the database would need to be considered at to determine whether it is related to another. This approach is inefficient, particularly as the size of a database grows, and the cost of traversing the entire store increases.

It is important to note that graph database systems are different to graph pro-

cessing frameworks such as Ligra [11] or Pregel [8]. Where a graph database system is concerned with fast and frequent storage and retrieval of graph entities, a graph processing system will typically perform more complex analysis and transformations of the graph, often in a lower volume. The distinction is similar to that between Online Transaction Processing (OLTP) and Online Analytical Processing (OLAP) systems. An example query for a graph database may be to find a route between two points in a graph, but an example query for a graph processing system might be to calculate pagerank for the network. This is an important distinction for our purposes, as the systems we discuss here are concerned mainly with data retrieval.

Some well-known graph databases include Titan, OrientDB and Neo4J. McColl et al. [12] provide a recent overview of these and others, as well as a performance comparison. In this paper, I restrict my attention to Neo4J. Neo4J is a schema-less database based on a key-value store, which runs on the JVM and presents a Java API.

## 2.3 Related Work

The incumbency of relational databases coupled with the potential scalability of graph databases naturally prompts the question of whether it is possible to efficiently store and query graph structured data in a relation database. Indeed, several efforts have been made in this direction.

Object-relational mapping (ORM) frameworks such as Hibernate ORM for Java allow programmers to map an object-oriented domain model to a relational database. This allows semantic links to be easily followed from one entity to another by abstracting away from the underlying relational stores. Although ORM frameworks allow programmers to feel as though they are using a graph database, they do not tend to focus on bringing any of the performance or scalability advantages of graph databases to the relational world.

Another approach is taken by OQGraph[**oqgraph**], which provides a plu-

gin for MariaDB and MySQL. By imitating a storage engine, users query a proxy table, which OQGraph interprets as graph traversal instructions. Although this somewhat improves the programmer experience by providing a more graph-centric interface, performance is not comparable to pure graph databases. Additionally, since the plugin imitates a storage engine, no parser extensions are possible and queries must be expressed in plain SQL. Although this means that no new query language need be learnt, it also means that graph-centric queries are awkward to express.

Biskup et al. [6] described in 1990 a SQL extension which could be used to query graph relations. No implementation is provided, and again, no performance considerations are made, but the work nonetheless provides a valuable inspiration for Grapht’s query language *gSQL*.

A more performance-oriented approach to bringing relational and graph databases together can be found in Sun et al.’s [13] SQLGraph. This approach uses relational storage for adjacency information, and JSON storage for vertex and edge attributes. Although this approach does improve performance for graph-centric queries, moving attribute information to JSON files hurts performance of normal relational queries, reducing performance in this direction.

Most recently, a Spark package has been produced which extends the *DataFrames* library to handle graph data called *GraphFrames* [**graphframes**]. DataFrames are a way of abstracting any underlying data source by loading it into a Spark *Resilient Distributed Dataset* (RDD). The GraphFrames extension extends these RDDs by allowing them to be interpreted as adjacency tables, and providing graph traversal methods. The query language provided is somewhat limited, however, primarily designed for simple pattern-matching within the graph rather than performing intelligent traversals as Grapht can do. Unlike Grapht, by building directly on top of Spark, GraphFrames queries are easy to distribute and parallelise across a number of machines.

Finally, my work also takes inspiration from Yoneki et al.’s *Crackle* [10], which improves graph performance using a relational store by maintaining a

limited graph cache in memory, which is periodically filled using graph-aware prefetchers to ensure topologically local vertices are preserved in the cache. Crackle did not consider relational queries either, however, nor did it attempt to present a declarative hybrid query interface in the way that Grapht does. We replicate and extend some of the Crackle results in Chapter 4.



# Chapter 3

## Implementation

In aiming to produce a hybrid system somehow equally well suited to both classes of problem, two approaches are possible. The first would be to create a novel system from scratch which fully commits to neither underlying storage strategy; following instead some middle path to achieve optimality. A second strategy – and the one employed here in the design of Grapht – is to extend the capabilities of one system to mitigate its weaknesses by taking inspiration from the other system. In particular, Grapht extends a durable relational data store with a graph-centric prefetched cache.

In this chapter, I first give a brief introduction to the different components involved in resolving a query through Grapht. Following that, I describe the two main components in more detail: the cache prefetcher and the query processor.

### 3.1 System Overview

As mentioned earlier, Grapht extends a relational data store – a graph query layer is “grafted” onto the underlying store. It is itself made up of two parts: a row-centric and a graph-centric handler. Both components communicate directly (and separately) with the underlying store via standard SQL queries.

This allows any SQL-compatible data-store to be used fairly indiscriminately. In addition to this, Grapht also includes a query processor, allowing simple queries to be expressed in a syntax similar to SQL, but harnessing the additional power of the graph-centric handler. Graph algorithms and queries vary widely, such that it would be impossible to efficiently represent the full diversity using a single query language. For this reason, the query processor is designed to be used as a possible alternative to using the Grapht API directly. (TODO: Diagram of main components)

The Grapht API has a relatively small interface. Method invocations to the API are mostly handled exclusively by one of the row-centric handler or the graph-centric handler, allowing each to be highly optimised for the task at hand. A few higher-level methods then allow client applications to combine the results of graph and row-centric queries. The row-centric handler performs relatively little work, relying instead on decades of optimisation made in relational databases to efficiently retrieve data by passing queries directly on to the underlying store.

Within the graph handler, data is laid out in memory to optimise graph-centric access patterns. The handler contains a cache, by design of limited size, so that Grapht can be used on systems even with small amounts of memory available. When a vertex is requested through the API which is not available in the cache, the vertex data is obtained from the relational database. At the same time, a prefetcher is used to populate the cache with adjacent vertices, so that subsequent requests from the query processor can hope to be fulfilled directly from the cache.

Finally, the query processor receives string queries from a client application and breaks them apart into row-centric subqueries targeting the relational database directly, and graph-centric subqueries targeting the graph handler. Client queries are written using an extension of SQL I call *gSQL* (as in *graph-SQL*) inspired by Biskup et al[6]. The extension is a strict superset of SQL, allowing pure-SQL queries to be passed on directly to the underlying store. Pure-graph queries end up being handled exclusively by the graph handler, and hybrid queries are broken apart and dispatched appropriately.



## 3.2 Row Handler

As mentioned above, the row-centric handler performs relatively little work. Instead, queries are passed unchanged to the underlying store, and the resulting relations are passed on. One optimisation is that queries to the row-handler are lazily evaluated by default. This allows queries to be transformed or dropped altogether when doing so would be advantageous. One situation in which this may arise is as part of a more complex API method call. The method may need to join together results from a graph query and a row-centric query. If the graph query returns an empty result set, the row query does not need to be evaluated - the joined result will itself be empty.

## 3.3 Graph Handler

The graph-centric handler is slightly more involved than the row-handler in that it contains a cache, in which vertices are looked up before resorting to the underlying store. Internally, it maintains an indexed list of vertices in memory. Each of these vertices contains a list of edges, allowing for quick traversal without needing to pass through an index again. If the full graph could be loaded into memory, these edges would contain a pointer to the target vertex directly to avoid indirection through an index in the same way. In practice however, we do not expect to be able to always contain the entire graph in memory, and instead the edges contain an index to the target vertex, which may not yet be loaded into the cache.

When a vertex is requested through the Grapht API, the cache index is first checked to try to fulfill the request. If possible, the vertex is returned directly. If this is not possible, the graph handler queries the underlying data store to obtain the vertex data. At the same time, a prefetcher is used to obtain data for related edges and vertices. The requested vertex data is immediately returned to the query processor, and the prefetched related data is added to a work queue. This work queue is processed in a

separate background process which is responsible for adding new items to the cache, and removing them when the cache exceeds the memory allowed by the current Java virtual machine (JVM). When the size of the cache does exceed the allowed size, vertices are removed from the index according to a least-recently used strategy. That is to say, each time a vertex is requested, the access is logged, when a vertex must be eliminated from the index, the vertex with the least recent access time is chosen. The idea behind this is that vertices which have been recently visited by the query processor are likely to be visited again soon in the future, since the traversal of the graph will remain local to the most recent nodes. Thus these vertices should be kept in the cache, in preference to vertices which were only visited a long time ago.

The performance of two different strategies (*lookahead* and *block*) for prefetching was analysed in detail by Yoneki et al. [10], but I have here replicated and extended these results with a third strategy.

### 3.3.1 Lookahead Prefetcher

The first strategy examined by Yoneki simply performed a limited depth-first traversal of the graph, starting from the requested node. This traversal was performed using a recursive “Common Table Expression” (CTE). These types of query allow a result set to be built up by recursively joining a partial result set with another relation. As the lookahead depth increases, the size of the partial result set grows exponentially, such that the cost of joining becomes itself exponentially larger. For this reason, Yoneki identified that after a certain depth, the cost of the prefetch procedure grows larger than the avoided overhead. I replicated this finding, which is further discussed in Chapter 4.

### 3.3.2 Block Prefetcher

The second strategy used by Yoneki et al. exploited the fact that there is often a strong correlation between vertex properties and graph locality. For example, in a road network, the latitude and longitude of junctions can be expected to correlate with graph locality, since nearby junctions are far more likely to be connected by a road than distant ones. Yoneki terms this “semantic” locality, and uses it to quickly prefetch large blocks of the graph by simply filtering on latitude and longitude. The cost of this procedure does not grow as fast as for the lookahead prefetcher - for a block of size  $b$ , we would expect the cost of fetching all junctions within the block to be proportional to the number of junctions within the block – i.e. the cost will be  $O(b^2)$ . Although this is advantageous from a performance perspective, it places an important limitation on the data – namely that there must exist some property which is tightly correlated with graph locality.

### 3.3.3 Iterative Lookahead Prefetcher

I examined a third possible strategy for this report, inspired by the first two. The lookahead prefetcher is most portable, since it can be indiscriminately applied to any graph. However, it does not scale as well, such that when the cache is large, and it would be advantageous to prefetch many vertices in one go, the performance cost is too high, and a block prefetcher becomes more effective. The source of this cost penalty is that many edges are unnecessarily traversed repeatedly. For example, after one level of recursion, edges one hop away from the source are traversed, and their destinations included within the CTE table. After five levels of recursion, we should only be examining edges which are precisely five hops away from the source. However the recursive CTE does not distinguish between the edges which were added at the previous level of recursion and those which were added at the start of the procedure (and which do not need re-examination).

We can improve performance, then, by only expanding the fringes of the

search – and even then only those fringes which have not been previously examined. To achieve this, I implemented a slightly more involved prefetcher, which maintains for each request a set of vertices explored this time. It then performs a number of iterations – one for each level of lookahead – requesting just those vertices which were added to the fringe during the previous iteration. Requests look like this:

```
SELECT *  
FROM edges  
JOIN vertices  
  ON vertices.id = edges.from_id  
WHERE edges.from_id IN (1286, 1373, 1141, ...)
```

These operations are very fast, since the `from_id` field of the edge relation is indexed such that retrieval is a constant time operation. This type of indexing is not possible for the recursive CTE table. The trade-off here is that many requests are made from the data store, such that the communication and query-parsing latency are incurred repeatedly. However, as is discussed in Chapter 4, experiments showed that in practice this latency was still smaller than the join cost of the CTE for large lookahead depths, and in fact this iterative lookahead approach allowed lookaheads of almost ten times the optimal depth of the original lookahead prefetcher without any significant loss of performance. Additionally, A\* search algorithms ran faster when using the iterative lookahead prefetcher, presumably because the prefetcher offered guaranteed graph locality, where the block prefetcher could only base decisions on a presumed correlation.

As a result of these findings, the iterative lookahead prefetcher was used as the default Grapht prefetcher. An optimal prefetch depth was not so clearly identifiable, and it is likely that this optimality will vary based on factors such as cache and graph size as well as the degree of interconnectedness of the graph (since a graph with many high-degree vertices will have a faster-growing fringe during exploration).

## 3.4 Grapht API

The Grapht API allows client applications to query the underlying data store, while taking advantage of the row and graph-centric handlers to optimise performance according to the expected access patterns. There is one API method targeting the row-centric handler directly: `getRelation(sql)`. This method takes a string parameter `sql` representing a standard SQL query, and returns an iterable list of rows representing the result relation. In order to target the graph-centric query, one must first initialise the graph cache with a particular source. This is done by using the `createGraph(vertices, vertexKey, edges, sourceKey, targetKey)` method, which takes four string arguments. The arguments `vertices` and `edges` correspond to SQL queries representing the edge and vertex relations respectively. Normally, these strings will simply be the names of the tables corresponding to edges and vertices in the underlying store, but it is also possible to pass in a SQL query which will dynamically create an edgelist. The remaining arguments, `vertexKey`, `sourceKey` and `targetKey` are used to determine which fields of the chosen relations correspond to vertex, edge source and edge target identifiers respectively.

Once the graph is initialised, two simple methods target the graph-centric handler directly, retrieving the relevant entities from the graph according to a given ID: `getVertex(id)`, `getEdge(id)`. The objects returned by these method calls have accessor methods of their own. Vertex objects have access to the `getEdgeList` method to obtain a list of outgoing edges, and Edge objects have access to a `getTargetVertex` method to obtain a reference to the vertex this edge leads to. As mentioned earlier, vertices are not guaranteed to be present in the graph cache when they are requested, as they may have been displaced by the LRU policy. A request to an absent vertex will trigger a fetch from the underlying store. On the other hand, edges will always be present cache as long as their source vertex is. Finally, both Edge and Vertex objects share a `getAttribute` method, used for retrieving arbitrary data stored alongside the edge. Between these five methods, it is possible

to traverse the entire graph efficiently, starting and ending at any arbitrary vertex, and collecting any desired data along the way.

Almost any algorithm can be efficiently expressed using this interface for graph traversal and retrieval of data. For convenience, a few further methods are provided, which use the methods described above for themselves. The most powerful of these is the `getPaths` method, which is used to discover paths branching out from some source vertex. The design of this method is closely tied to the design of the query processor (discussed in the next section), since this method provides the bulk of the implementation for the query processor. However, a client application can also call `getPaths` independently.

The operation of `getPaths` essentially follows a best-first exploration of the graph starting from the source vertex. To achieve this, a simple priority queue of partially-discovered paths is used, where the priority for new vertices is zero by default (in which case the priority queue devolves into a normal first-in, first-out queue and the search proceeds as a breadth-first search). The priority of a vertex is otherwise determined by a user-provided *prioritiser* function.

Rather than provide a list of edges and vertices as output, a collection of *accumulators* are provided to `getPaths`. These are in effect stateful functions, which will receive new data at each step of traversal, and update their internal state accordingly. For example, a `count` accumulator could be provided, which simply increments its state with each new vertex encountered. Alternatively, a `sum` accumulator could be provided which will extract some property from the edges traversed, and add the values together.

Along with these accumulators, *evaluator* functions can be provided, which determine whether the state of the accumulator is such that the current path should be saved for inclusion in the result set. This allows us to, for example, limit traversal to a certain depth by using a `count` accumulator in conjunction with an appropriate evaluator. An opportunity to cut the search space down is also provided to the user, who can optionally specify

a direction or rate of change for each accumulator. For example, a `count` accumulator is guaranteed to increase by one with each step through the graph, while a `sum` accumulator can be assumed to at least increase each step if the property it is accumulating is always positive. If an evaluator is provided which rejects values greater than some given constant, then we do not need to continue expanding paths once that constant has been reached with the `sum` accumulator.

Two boolean values may also be provided, specifying that only unique edges or vertices be used along the path. This is determined by maintaining a “closed” set of already- visited entities for each path. If a vertex or edge which is already in a closed set is expanded, it is skipped rather than being added to the queue as a new partial path. Finally, a *limit* parameter is used to determine how many valid paths should be retrieved before returning. By default, the limit is unset, and the search for valid paths will be exhaustive.

To summarise, then, `getPaths` receives as mandatory input a graph and start vertex, and the following optional parameters:

- A *prioritiser* function to guide traversal
- A list of vertex *accumulators* (optionally with a rate of change specified)
- A list of edge *accumulators* (optionally with a rate of change specified)
- A list of *evaluators* to determine valid paths
- Two boolean values specifying whether vertices and edges must be unique along a path
- An integer *limit* on the number of paths retrieved

In the special case where a single path is requested, an important optimisation is possible. When multiple partial paths pass through the same vertex, we need only consider the route with the highest priority. In highly interconnected graphs, this allows us to discard a large number of possible paths early on. In practice by doing this, and requesting only unique vertices, this optimisation allows the best-first traversal to degenerate into a variation on

the A\* search algorithm, as long as an admissible heuristic is used in the prioritiser function to guide traversal.

On each iteration of the traversal, a partial path is fetched from the priority queue. This path is considered to see if it satisfies the conditions required to be output. If it does, the path is saved, and we check whether the desired number of paths have been found, stopping execution if this is the case. If not, we consider all possible extensions of this path by iterating through the outgoing edges of the final vertex in the path. If it is possible that these edges may lead to a valid solution (which can be determined by examining the attributes' rates of change, as discussed earlier), then the target vertices are fetched from the cache. A new partial path is created, accumulating attributes from both the edge and the target vertex. A priority is calculated for this path according to the prioritiser function, and the new path is added to the queue, ready for selection in some future iteration.

The eventual output of `getPaths` is a list, where each entry contains the accumulated results of traversing a valid path from the source vertex. In effect, this result can be treated as a relational table, where each accumulator is a field in the relation. This allows paths through the graph to be manipulated just like any relation in the underlying store, and in particular allows the results of a path-finding query to be joined with a normal relation. This ability is a significant contribution of my work, since obtaining this kind of result efficiently is not possible in either a purely relational system (where the path-finding component would take too long), nor in a purely graph-oriented system (where there exists no notion of a normal relation). Ideally, the result relation of the `getPaths` method would be created directly within the underlying store, so that joins could literally be performed as between normal relations. The current Grapht prototype is not tightly coupled enough to the underlying store for this to be possible, however. As a workaround in the meantime, the results of the `getPaths` method can be loaded into the store after they are calculated as a temporary in-memory relation. This does carry an inevitable performance penalty compared to being constructed directly in PostgreSQL's memory space. The precise impact of this prototype limitation



is discussed further in Chapter 4.

It may seem as though only providing a path-searching method does not provide a very powerful abstraction beyond the vertex and edge-querying methods provided. However, the aggregation mechanism means that we are able to discard path information if it is not of interest, and instead only consider data belonging to nodes in some way reachable from the source node. For example, we may wish to retrieve a list of all junctions reachable within 100km of some given source junction. This can be achieved by using a `sum` accumulator as described above and a `last` accumulator, which simply retrieves a certain property from the last vertex along a path. An evaluator can then be attached to the `sum` accumulator, checking whether the path to reach this vertex is less than 100km. Since the sum will be non-decreasing, an exhaustive search will terminate as paths cease to be explored when 100km are passed. In this way, we can quickly retrieve a relation containing just the details of nearby junctions using a single API method call.

## 3.5 Query Processor

The aim of the query processor is to provide some of the power of the Grapht API in a more lightweight manner, so that clients can simply express retrieval queries as strings, rather than needing to interact with a more low-level API. Biskup’s Graph Manipulation Kernel[6] (GMK) unites relational and graph models in a way similar to what we would like to achieve. It does this by defining paths over graphs as relations with three implicit attributes: `START`, `GOAL` and `LENGTH`. By describing paths in this way, we write queries over the possible search space of all paths through the graph using a syntax very similar to SQL. For example, a “friend-of- friend”-style query to find vertices two hops away from some the vertex with ID 1 would be written in Biskup’s graph manipulation kernel:

```
SELECT *  
FROM PATHS OVER my_edge_relation
```

WHERE

START = 1 AND

LENGTH = 2

In this example, `my_edge_relation` is some relation with precisely two attributes which are interpreted as start and end identifiers of edges in a graph. If the relation contains more than two attributes, the attributes indicating start and end may be included in parentheses, as in “... PATHS OVER `my_edge_relation` (`id1`, `id2`).” Other properties of the edge relation can then be aggregated along the path to provide new data to be included in the output path relation. Although this provides a good starting point for Grapht, a few notable shortcomings prevent us from adopting the solution directly.

Firstly, Biskup’s solution has no direct support for vertex attributes. Vertices are in fact not treated as first-class entities at all, simply being included implicitly as start and end properties within the edge relation. It is possible to include vertex properties in a query by joining the edge relation together with the vertex relation, however this results in duplication of vertex data when several edges lead to the same vertex, and can cause confusing “off-by-one” errors since any path will have one more vertex than it has edges.

Secondly, Biskup’s extension leaves the traversal mechanism used to find valid paths unspecified. In many ways this is a good thing, and mimics the design of SQL itself. By ensuring that the query is fully declarative, an implementation can select any method to find possible paths, and the client application need not worry about efficiently finding solutions. In the absence of cycles in the graph, the set of paths discovered will be constant, although they may be discovered in a different order. This is no different to the way in which the order of rows retrieved by a SQL query is undefined unless an `ORDER BY` clause is present in the query. The difference is that the performance gain from using an appropriate traversal order through the graph is much more significant than the gains made by iterating through a set of rows in a different order. If we want to find a particular neighbour of a vertex, a depth-first search (DFS) may visit every single other node in

the graph before returning, while a breadth-first search (BFS) can terminate almost immediately.

Adding cycles to the graph makes appropriate selection of a traversal order all the more critical, since DFS here may result in non-termination, while BFS may have no problems. A final related problem with GMK is that it is difficult to express uniqueness constraints on the vertices visited. By enforcing that the vertices visited along a path be unique, it becomes possible to traverse a graph with cycles without non-termination, so expressing these constraints is very valuable.

### 3.5.1 Improving on the Graph Manipulation Kernel

To solve the issue of vertex attributes, two small changes must be made. Firstly, defining a graph according to an edge relation is no longer sufficient, and graphs must now be specified as a pair consisting of an edge relation and a vertex relation. As before, we must specify the attributes to use as start and end points for the edge relation, or vertex identifier for the vertex relation. We must also now change the syntax of calculated attributes of the output relation, to explicitly state whether values are being accumulated along the edges of the path, or along the vertices. For example, to accumulate the `name` property of vertices as a concatenated string, along with the total `cost` of edges along the path, a gSQL query may look like this:

```
SELECT START, END, LENGTH,
       (ACC VERTICES CONCAT(name, " -> ")) path,
       (ACC EDGES SUM(0, cost)) cost
FROM PATHS OVER (my_edge_relation(id1, id2), my_vertex_relation(id))
WHERE
    START = 1
```

In this way, the accumulation is made explicitly over vertices or edges (`ACC VERTICES` in the former case, `ACC EDGES` in the latter). Following this declaration of accumulation is a function which will be used as an accumulator

(**CONCAT** or **SUM** in the example). For the prototype version of Grapht, only a few functions were defined but others would be trivial to add. The **CONCAT** function, for example, takes as arguments the name of a property to accumulate at each step and a separator string, returning the concatenated value of all named properties along the path.

To solve the remaining problems, I have added a **TRAVERSE** clause to the end of the query format, which allows users to specify a prioritisation order for traversing new nodes in search of a valid path. The clause is optional, and if omitted a breadth-first search is used. An example use of this clause would be:

```
SELECT *
FROM PATHS OVER (my_edge_relation(id1, id2), my_vertex_relation(id))
WHERE
    START = 1
TRAVERSE
    UNIQUE VERTICES
    BY MIN(length)
```

This example specifies a breadth-first traversal, avoiding cycles by disallowing repeated vertices along the same path. The **UNIQUE** modifier is again optional, and may specify either “**VERTICES**”, “**EDGES**” or “**VERTICES,EDGES**” if duplicate edges and vertices should both be avoided. **length** is a default property of all paths, and so at each opportunity to discover a new vertex, Grapht will select one according to the minimum path length so far (i.e. according to a breadth-first traversal).

One limitation has been introduced to gSQL queries which was not present in GMK. This is that queries should always specify a start vertex for the search. This is included for performance reasons, following the implementation of the **getPaths** method of the Grapht API. In general it would be inefficient to consider every possible path through the graph when the query can be solved using a local traversal. It would be possible to reduce the search space without using a known start vertex (for example by working backwards from

a known end node) but these strategies have not been implemented for this initial prototype, hence a start node must always be provided.

### 3.5.2 Executing gSQL

The main task of the query processor is to extract information from a gSQL query in order to call the `getPaths` method of the API. When a gSQL query is received, the first task is to separate components of the query into row-centric and graph-centric queries. To achieve this, the source of data specified in the query is looked at. If data is being queried from one of the relational tables, then that data is fetched through the row handler directly. If data is being queried from a “`PATHS OVER`” clause, then the query will instead be resolved through the graph handler. This provides a very simple way to break apart a complex query and reliably identify what kind of access pattern should be optimised for.

In the case of complex clauses which join together results from disparate data sources, a relational join can be performed as between normal relational tables. This is possible since the result of a “`PATHS OVER`” query (i.e. the output of a call to `getPaths`) When a query has been identified as targeting the graph handler, it is further broken down into actionable components. First, the data source is identified from the `PATHS OVER` clause. If the precise combination of edge and data relations has not been seen before, a new subgraph is initialised in the cache to represent it. Next, accumulator functions are generated to collect the data required by the `SELECT` clause. Any accumulators specified in the query are created according to the parameters passed in the query. In addition to these, three accumulators are always created to represent the `START`, `LENGTH` and `END` attributes of the output relation. Finally, accumulators may also be defined within the `WHERE` clause, so that conditions can be checked at each step even if the result is not included in the final relation. This `WHERE` condition is evaluated at each vertex along a possible path.

The `LENGTH` accumulator which is implicitly created is initialised with a known

rate-of change. This allows queries which have a bounded depth in the **WHERE** clause to run exhaustively within the appropriate bounds, since Grapht knows that a shorter **LENGTH** will never arise further along the path. Although this is a desirable property, there is currently no way for query writers to specify that other accumulated variables have a known rate of change, although future versions of the query processor may provide functionality for this.

The prioritiser and uniqueness constraints for **getPaths** are very simply extracted from the **TRAVERSE BY** clause. Similarly, a limit to the number of paths to retrieve is read directly from the **LIMIT** clause of the query.

# Chapter 4

## Evaluation

In this chapter, I provide experimental evidence for some of the claims made in earlier chapters. I will begin by describing my experimental setup, and outline the methods used for subsequent tests. Three sections follow, each examining a different performance aspect. The first of these sections aims to replicate extend the results of Yoneki et al[10] in determining optimal an prefetching strategy for Grapht. The next section examines the performance of certain queries across the three database engines under consideration: PostgreSQL, Grapht and Neo4J. This allows me both to confirm the premise motivating the development of Grapht – that there exists a large performance gap in both directions between PostgreSQL and Neo4J – and to show that the addition of an intermediate hybrid query processor can help to bridge this gap and improve average overall performance. Finally, the last section examines includes a more qualitative examination of the expressiveness of the three query systems, both in terms in terms of the declarative query language presented, and in terms the direct API offered (where one is available).

## 4.1 Test Framework

Before undertaking any tests, an experimental hypothesis was chosen to direct the direction of my investigations. These hypotheses are outlined in the appropriate section. To obtain a measure for query performance, a Scala benchmarking program was written, which recorded the start and end time of operations and used these to calculate the average execution time for each. Each query was performed ten times, with the fastest and slowest results discarded as outliers. Unless otherwise stated, 100 queries of approximately the same size were randomly chosen for each test, to prevent the particularities of any single query from impacting the overall results. The execution time listed is the mean time taken across all repetitions of all queries. For PostgreSQL and Neo4J, a cache-warming query was performed before each test by iterating through every vertex in the graph. This was not performed for GraphT, since such a cache-warming query would cause the prefetcher to load much of the graph into memory, making the presence of the prefetcher redundant by the time the query is performed. This disparity is tolerable, since we nonetheless will see a conservative estimate for GraphT’s performance with respect to our performance hypotheses.

The same dataset was pre-loaded into both PostgreSQL and Neo4J, consisting of a DIMACS dataset representing the U.S. road network[**dimacs**]. This graph contained XXX vertices and YYY edges. Both edges and vertices were then augmented with a “**payload**” attribute calculated as the MD5 hash of the entity ID. This was done to ensure that the database contained more than simple topological information, as is likely to be the case for many practical applications of graph databases. In PostgreSQL, the graph was defined by two relations, as shown in Table ???. In Neo4J, the graph was defined according to the standard graph format, with an index created on vertex IDs, and attributes added to vertices and edges as entity *properties*.

All experiments were carried out using Ubuntu 16.04 running on a 64-bit X-core Intel Core i7 processor (Yth generation), with Z GB of RAM. JVM vQQQ was used, with the maximum and minimum heap size set to 1024MB



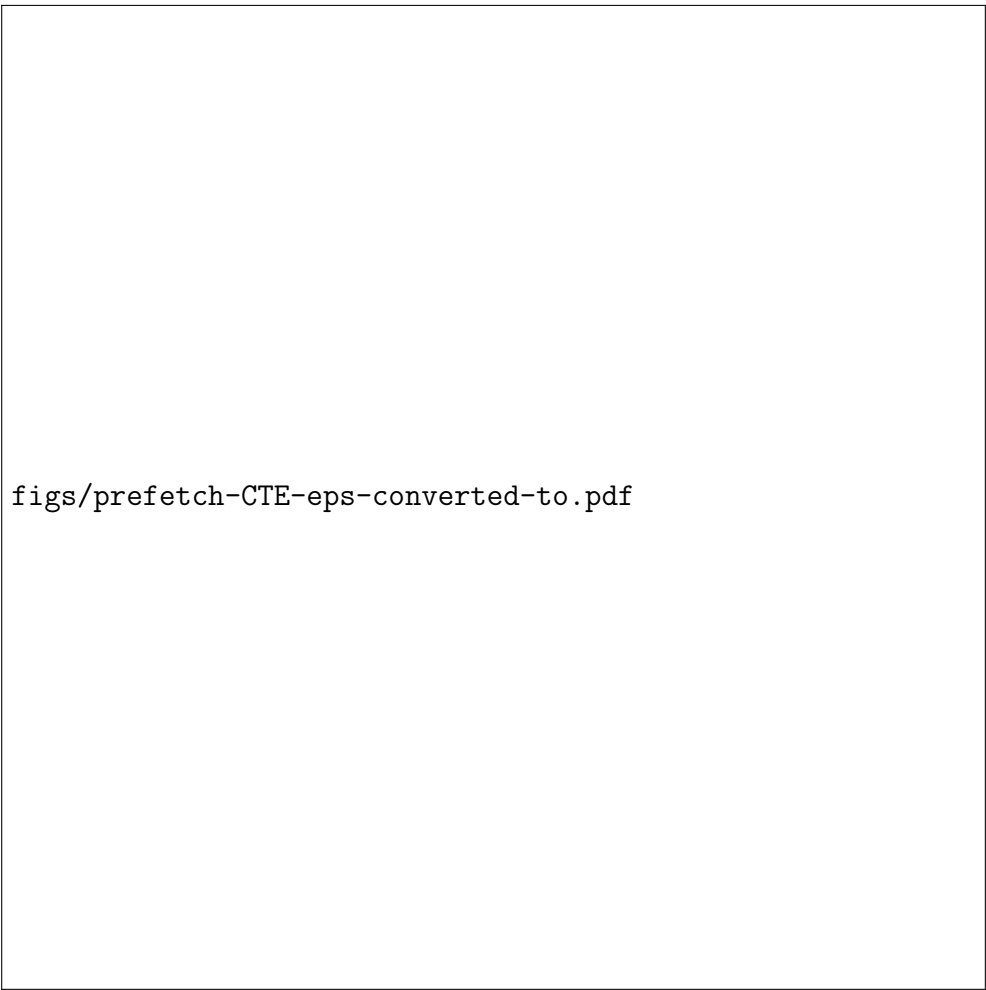
by default, to avoid wasting time trying to acquire more memory from the system. The JVM has global pauses for garbage collection, so a manual garbage collection was requested before each test, to try and ensure that any pauses were at fault of the current test itself.

## 4.2 Prefetcher Performance

The first task was to select a prefetching strategy to use in Grapht. For this evaluation, I examined the prefetchers described in section 3.3. In evaluating prefetcher performance, the initial expectation was to find that the peak performance would come from prefetching with large blocks rather than using a lookahead strategy – as reported by Yoneki et al. In addition, where the original CTE prefetcher suffered as lookahead depth increased, I expected to find that my iterative lookahead prefetcher would perform better by performing less expensive queries.


To test the effects of changing prefetching strategies, the execution time for an A\* search through the prefetcher was measured. I consider here separately five different query lengths, in order to see whether prefetching strategies were particularly good or bad when examining a larger portion of the graph. Paths consisting of 20, 40, 60, 80 and 100 hops were each considered. Twenty queries were found for each path length by randomly selecting start and end points within the same graph neighbourhood, and performing a shortest-path calculation to select only paths of the desired length.

Figure 4.1 shows an analysis of the effect of changing the depth of the CTE lookahead prefetcher. As expected, we see that for all path lengths performance initially improves as we increase path length, until a lookahead depth of around 4 to 5. At this point, execution time begins to increase again. We can gain a better causal understanding of this performance cost by examining the number of vertices needing to be fetched from the database for each query (the graph *miss rate*). We see from Figure 4.4 that the miss rate decreases exponentially as lookahead depth increases. This is expected since



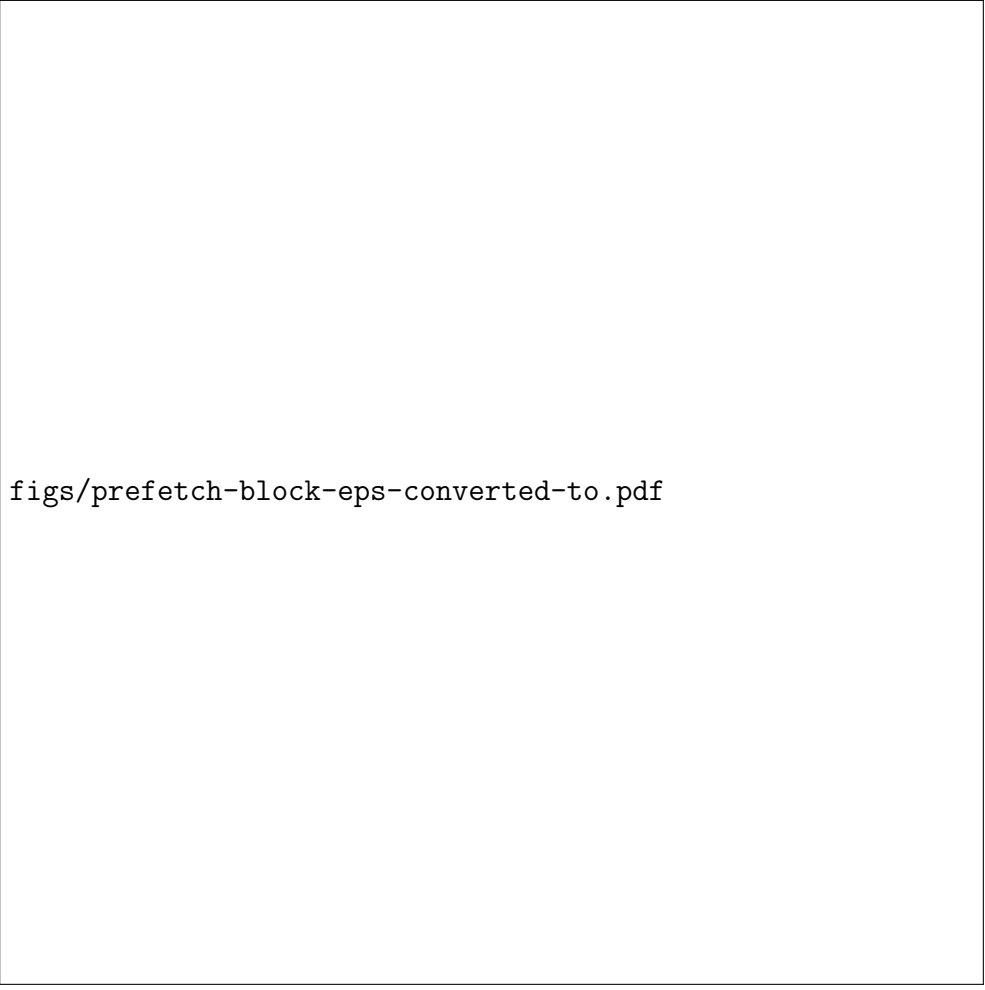
figs/prefetch-CTE-eps-converted-to.pdf

Figure 4.1: Average Execution time for A\* search of different path lengths using a CTE lookahead prefetcher



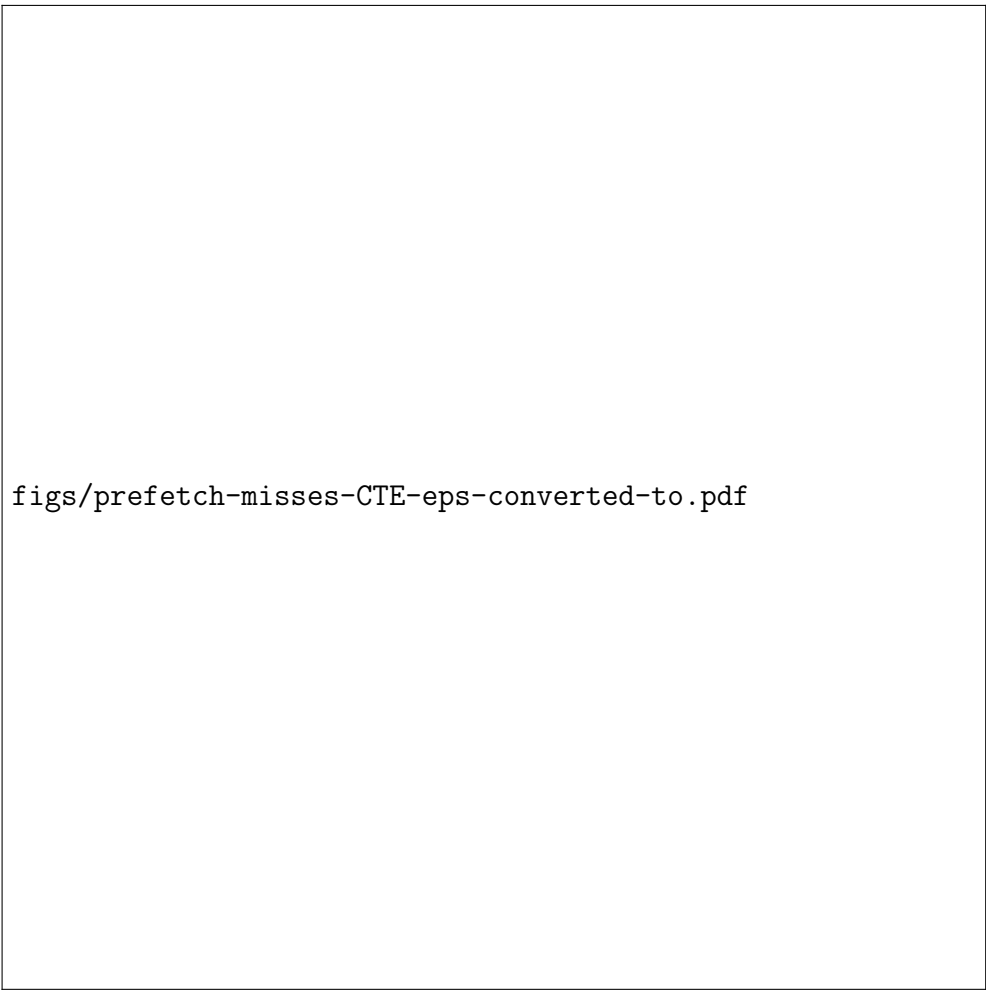
figs/prefetch-itr-eps-converted-to.pdf

Figure 4.2: Average Execution time for A\* search of different path lengths using an iterative lookahead prefetcher



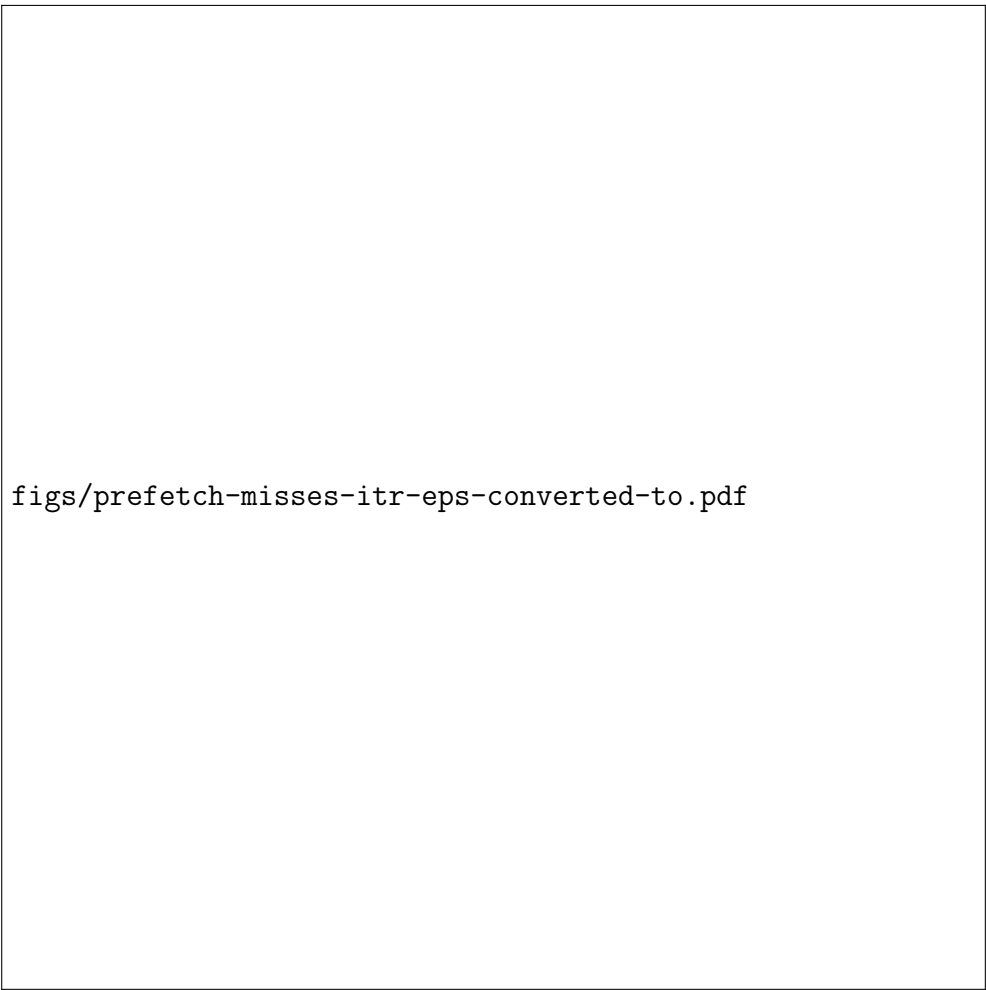
figs/prefetch-block-eps-converted-to.pdf

Figure 4.3: Average Execution time for A\* search of different path lengths using a block prefetcher



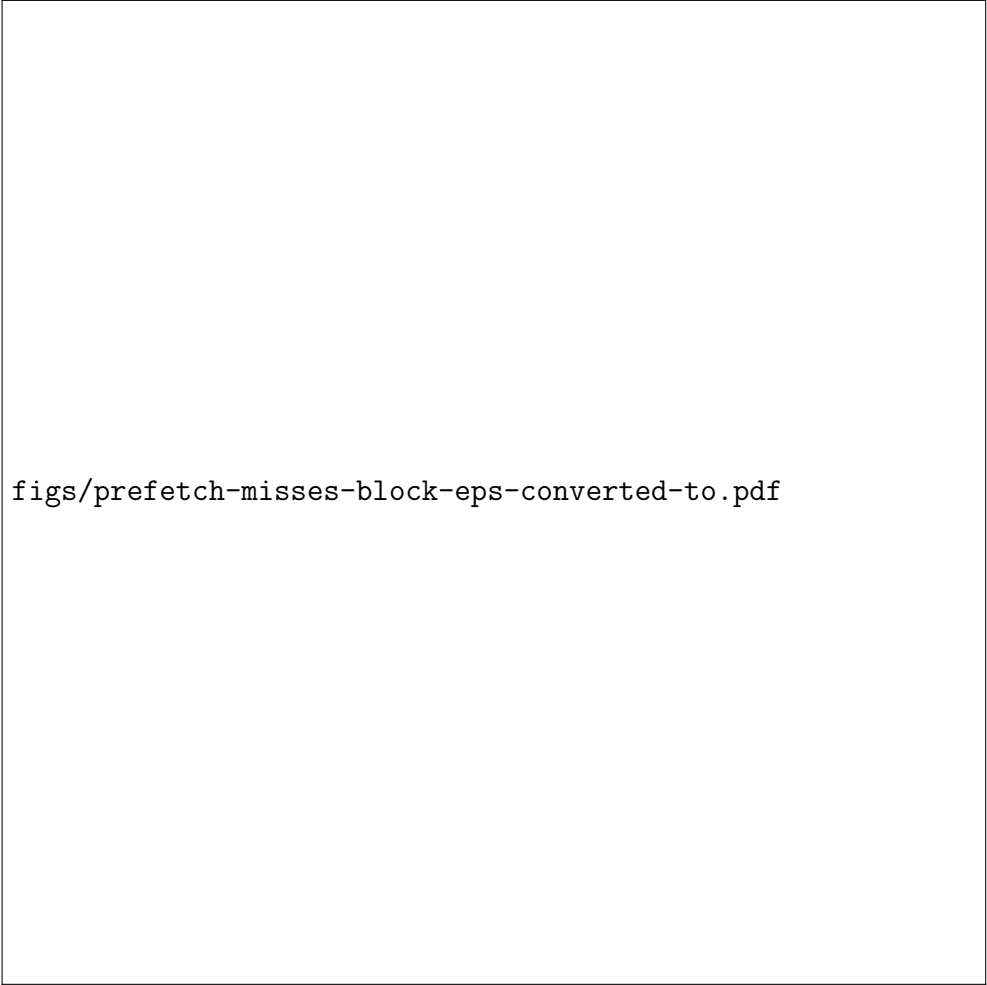
figs/prefetch-misses-CTE-eps-converted-to.pdf

Figure 4.4: Graph miss rate for A\* search of different path lengths using a CTE lookahead prefetcher



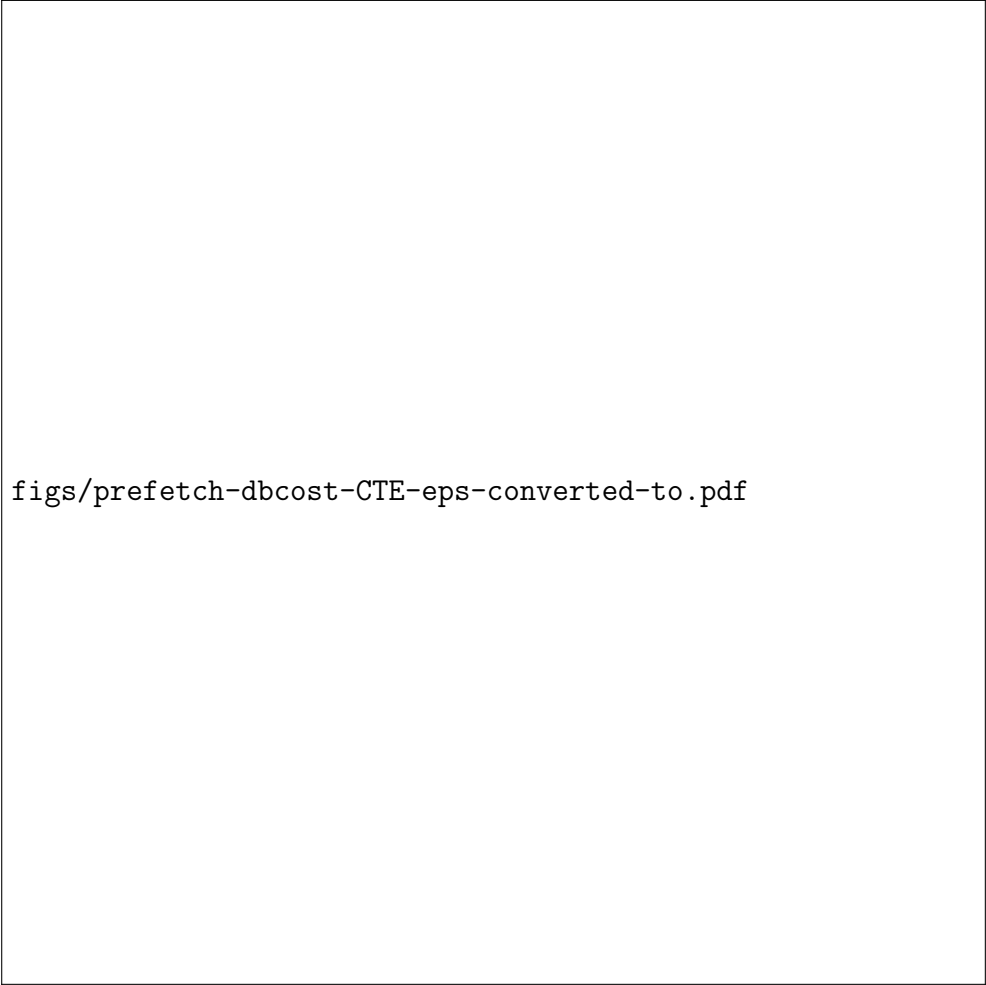
figs/prefetch-misses-itr-eps-converted-to.pdf

Figure 4.5: Graph miss rate for A\* search of different path lengths using an iterative lookahead prefetcher



figs/prefetch-misses-block-eps-converted-to.pdf

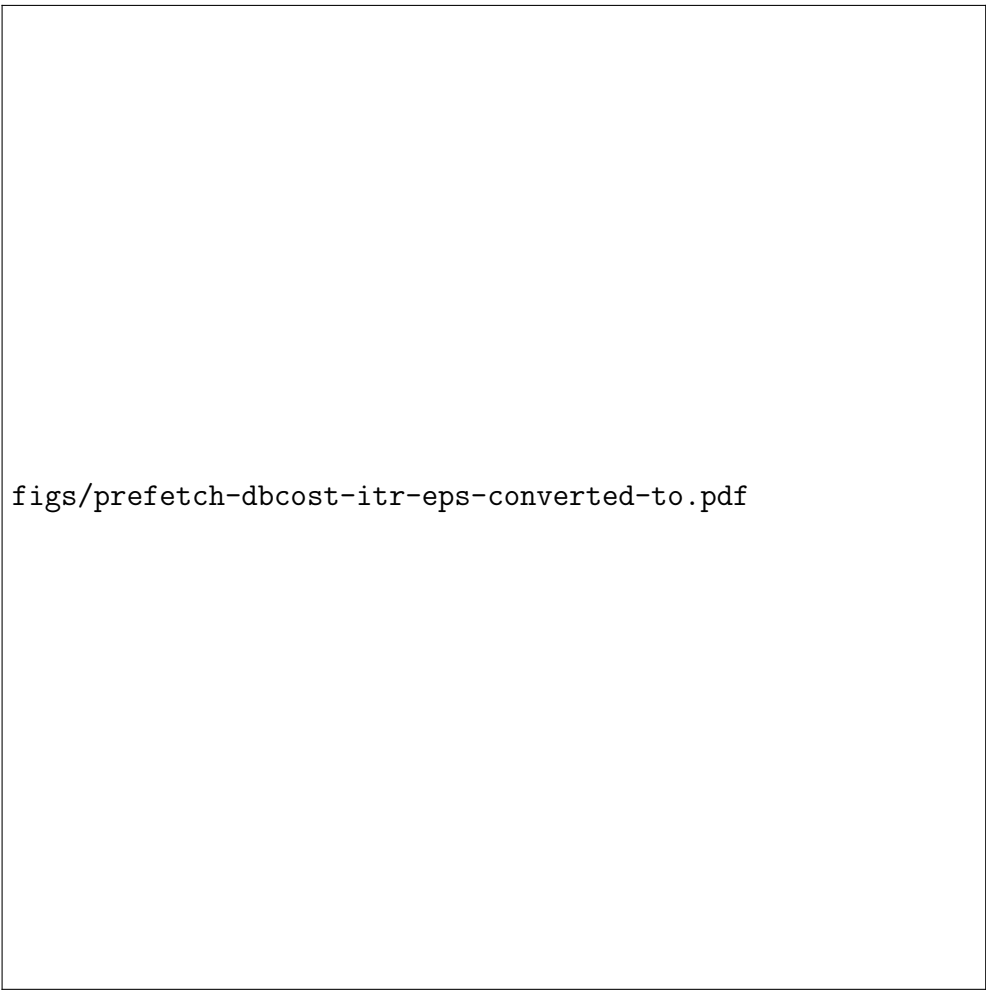
Figure 4.6: Graph miss rate for A\* search of different path lengths using a block prefetcher



figs/prefetch-dbcost-CTE-eps-converted-to.pdf

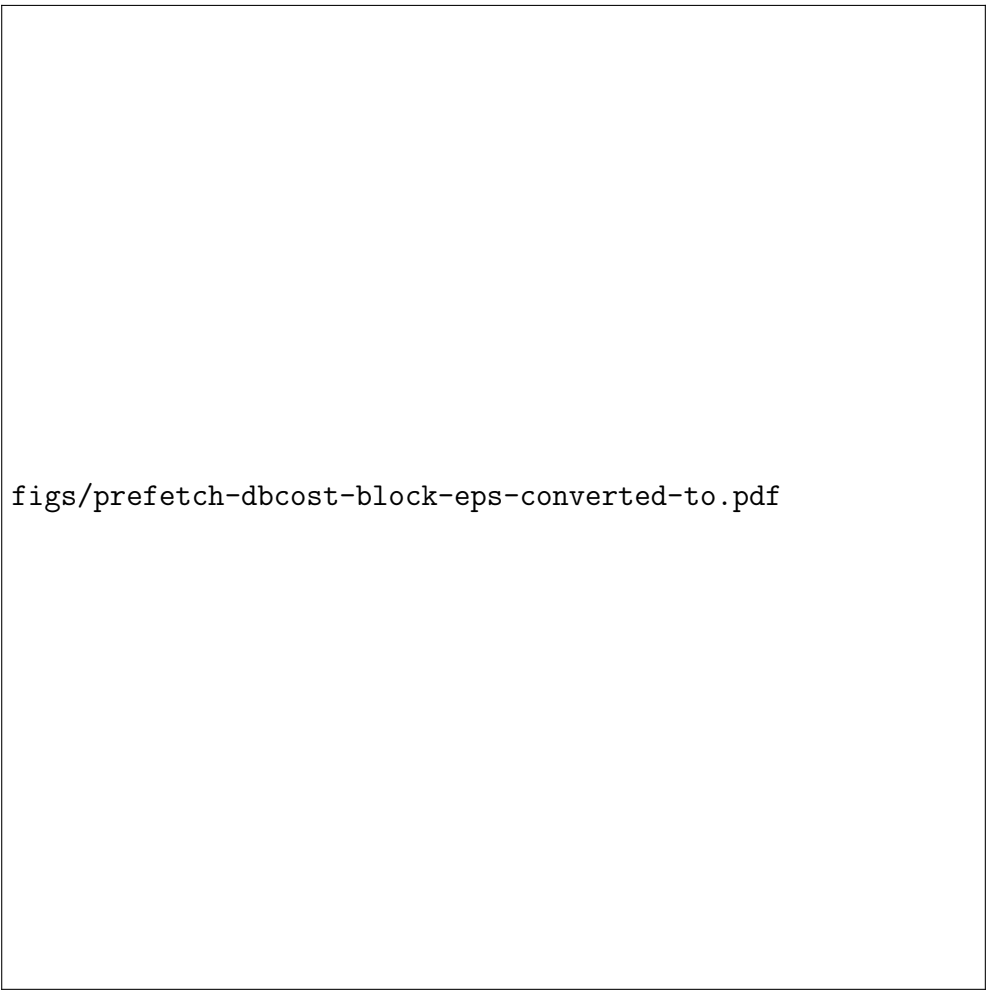
Figure 4.7: Average cost of cache miss for A\* search of different path lengths using a CTE lookahead prefetcher





figs/prefetch-dbcost-itr-eps-converted-to.pdf

Figure 4.8: Average cost of cache miss for A\* search of different path lengths using an iterative lookahead prefetcher



figs/prefetch-dbcost-block-eps-converted-to.pdf

Figure 4.9: Average cost of cache miss for A\* search of different path lengths using a block prefetcher

more vertices are fetched with each miss. The overall execution time thus initially decreases with depth. In contrast, the average cost of each cache miss increases super- exponentially (Figure 4.7). This cost quickly overcomes the benefit from the reduced number of graph misses, and becomes the dominating factor as lookahead depth increases.

For the iterative lookahead strategy, we see that the average cost of each graph miss (Figure 4.8) is initially higher than for the CTE strategy. This is due to the fact that several queries are performed at each iteration instead of one more complex query. However the average cost using the iterative approach grows far more slowly with lookahead depth, allowing lookahead depths around five times larger than the CTE approach to be used before performance starts to degrade. This larger lookahead depth in turn causes a much lower graph miss rate. Indeed, in all of the examined queries, the lookahead depth was allowed to exceed the path size, meaning that the cache miss rate effectively settled at one (when only the source vertex misses). The effect of these two factors on overall database time is that the total time taken fetching vertices from the database decreases steadily as the cache miss rate decreases. As soon as the graph miss rate reaches 1, however, the overall time taken increases with lookahead depth, since additional iterations only fetch vertices which will not be examined by this query. Although the peak performance of this approach is similar to the CTE lookahead strategy for A\* search, the overall performance characteristics are quite different. It seems likely that the vastly increased lookahead depth would improve performance for other queries where the graph miss rate is never allowed to reach 1.

Finally, we examine the effects on performance of using the *block* prefetcher (Figure 4.3). Predictably, the overall execution time of A\* reduces as the block size increases. This happens because the average cost of a graph miss remains roughly constant as block size increases (Figure 4.9), but more vertices are added to the graph each time. This reduces the graph miss rate (Figure 4.6), which causes the decrease in execution time. However, I found that this reduction slows after a block size of around 15000, with an overall execution time of between one and ten seconds. Since both lookahead

prefetchers achieved execution times of around 0.1s, and the rate of improvement for the block prefetcher had stalled, no further experiments were run at larger block sizes. It is interesting that this result contradicts the earlier one from Yoneki et al., who found that the block prefetcher offered better performance than a lookahead approach. It is possible that this disparity is due to improvements to PostgreSQL in the meantime, or down to mis-configured indexes in the edge relation. These indexes are critical to achieve good lookahead performance, and indeed without them a block prefetcher performs significantly better.

For subsequent experiments, I used a CTE lookahead prefetcher with lookahead depth 5, since this offered a reliably good performance across all path lengths.

### 4.3 Query Performance

Having established an appropriate prefetching strategy, the next task is to verify the motivating hypothesis – that there exists a significant performance difference between PostgreSQL and Neo4J for certain queries. Informally, we expect Neo4J to perform better for “graph-centric” queries, whereas we expect PostgreSQL to perform better for “row-centric” queries. A graph-centric query will tend to have the following features:

- The absence or existence of edges between vertices is of greater interest than the properties of the edges or vertices.
- Only a local part of the graph need be considered to satisfy the query

In contrast, a row-centric query will place equal weight on the properties of entities as on the topology of the graph, and will typically need to consider vertices and edges which are not necessarily connected to one another.

Specific queries were chosen as representative of each of these categories, motivated by the mapping problem described in Chapter 1. To represent graph-centric queries, an A\* search query was used. This satisfies the requirements,

since the presence or absence of edges is what characterises a path between two points, and properties of the vertices are otherwise ignored. Additionally, a solution may be found by searching through only nodes local to the start point, rather than needing to look at the entire graph. To represent row-centric queries, a string-matching query was chosen, which identifies the node with a payload which most closely matches some user-specified string. A match is determined using by calculating the Levenshtein distances<sup>1</sup> between the two strings. This is close to the kind of query required to find a location based on an approximate string search in a mapping application, and is clearly row-centric since the vertex payload is of main interest, and every vertex in the graph must be considered.

### 4.3.1 Performance of Graph-centric Queries

To generate random queries for the A\* search, a similar methodology was employed as for the prefetcher queries described in Section 4.2. Paths were chosen to be of specific lengths between 20 and 100 hops, so that the average solution time would not be dominated by the execution time required to find particularly long paths. Twenty routes for each path length were examined.

Aiming to minimise variation in performance due to differences in three implementations of A\*, a single Scala implementation was prepared, with three adaptors responsible for fetching data from the relevant store. For Neo4J and Grapht, the direct API was used rather than the declarative interface. Although it is possible to create an implementation of the A\* algorithm entirely in SQL, the lack of an efficient way to create a priority queue makes performance unfairly bad. Similarly, the Neo4J API provides a direct `astar` method which performs the full computation internally, but it would be unfair to compare its performance against another system which needs to transfer data from store to external application at each iteration, since Neo4J here does not suffer from the transfer latency. By providing a single A\* implemen-

---

<sup>1</sup>Informally, the Levenshtein distance is the minimum number of single-character edits required to change one word into another

tation for all three databases, we reduce our test to focus solely evaluating their performance as a data store, rather than as an implementation platform.

Figure 4.10 shows the average time taken to find each path. Note that a logarithmic scale has been used for clarity, since there is a performance difference of several orders of magnitude between the systems. We first notice that Neo4J does indeed provide a very significant performance gain over PostgreSQL for this type of query of nearly 1000x across all path lengths. This is unsurprising since A\* is an archetypal example of a very local, topology-focussed algorithm, around which we expect a graph-centric database system to be optimised.

As hoped, Grapht bridges the performance gap between the two systems. Although Neo4J’s optimisations still grant it a performance advantage of around 35x, Grapht completes search queries almost 30x faster than PostgreSQL on average. This is a particularly pleasing result given that this performance improvement has been given “for free”, in that no changes have been made to the underlying store at all. By simply redirecting queries to pass through Grapht instead of directly to PostgreSQL, A\* performance can be significantly improved.

### 4.3.2 Performance of Row-Centric Queries

To generate Levenshtein distance queries, a random number was generated, and an MD5 string was created from it. This string was used as the user-supplied search string. Figure 4.11 shows the time taken for each of the database systems to identify the vertex with the smallest Levenshtein distance from the user-supplied string. Here, we see the opposite results to the A\* query. PostgreSQL completes the query significantly faster than Neo4J, by about 6x. This is not such a large performance gap as the 1000x seen in the previous subsection, it seems likely that the cost of the query is largely dominated by the cost of performing the Levenshtein comparison. This cost will be comparable for the two engines, with the performance difference instead coming from PostgreSQL’s row-centric optimisations.

Figure 4.10: Average A\* Execution time for paths of varying length across three database engines

Figure 4.11: Average Levenshtein-search execution time for 100 queries across three database engines



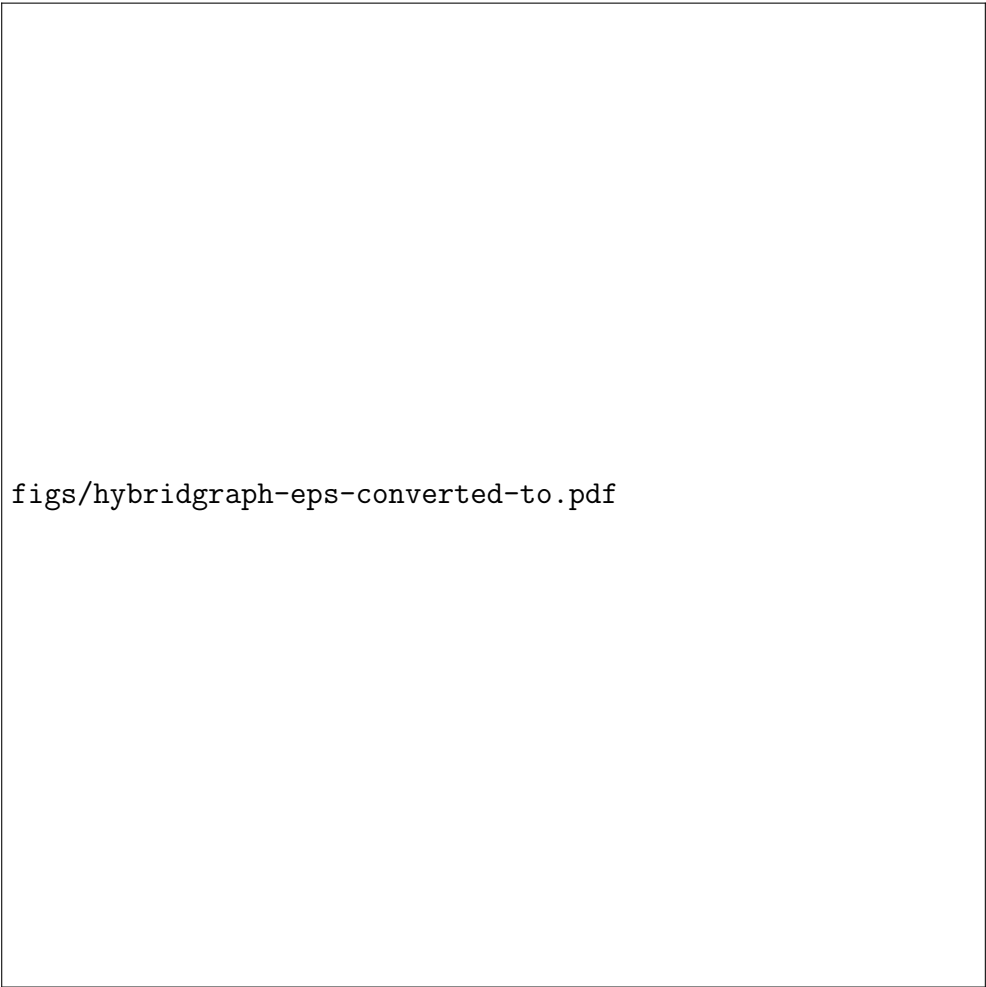
Since Grapht builds on top of a standard relational database, it is able to simply forward the query on to PostgreSQL and achieve exactly the same performance. By doing this, we have improved significantly on Neo4J’s worst-case performance.

### 4.3.3 Performance of Hybrid Queries

Finally, we wish to consider not simply the performance of graph and row-centric queries in isolation, but also when they are combined as in our motivating example. To create this kind of query, we use first a Levenshtein distance query to select a target vertex, and then perform an A\* search from some pre-selected source vertex (as if our user was searching for a route to some destination by name). An appropriate source vertex was selected before tests were run in order to ensure that the shortest paths were of a comparable length of 50-60 hops.

Figure 4.12 shows the average execution time for these hybrid queries. In this scenario, Neo4J once again maintains a lead over PostgreSQL, thanks to it’s 1000x performance difference for A\*. More interesting, however, is the fact that Grapht, by improving upon the worst-case performances of both other systems, has managed to outperform both commercial systems.

In reality, it would be impossible to select a database system based on these benchmarks for use in some given scenario, without first determining exactly the type of queries which will most often be encountered. What these results do tell us, however, is that for workloads fairly evenly balanced between graph and row-centric queries, a hybrid system is able to outperform both alternatives. If a workload is identified which is dominated by graph-centric, it is likely that Neo4J will not suffer enough from the occasional row-centric query to outweigh the benefits of the graph-centric optimisations. However if the workload is largely row-centric, Grapht will generally perform better. Note also that Grapht will almost always perform either equally or better than PostgreSQL alone, suggesting that an in-memory graph layer may be a worthwhile improvement to make more closely within the relational system,



figs/hybridgraph-eps-converted-to.pdf

Figure 4.12: Average Execution time for 100 queries consisting of a Min. Levenshtein-distance search and an A\* path search across three database engines.

rather than needing to exist as an optional add-on.

#### 4.3.4 Expressiveness of Query Systems

In the previous subsections, I performed a quantitative analysis of the performance of three database systems for graph-centric, row-centric and hybrid queries. However performance is only one aspect contributing to the success of a system. Ease of use during development is another crucial factor. If new custom queries are written regularly to explore a dataset, then slow-running but quick-to-write queries may obtain an answer to a user query more quickly overall.

There are five interfaces to consider in this evaluation. Neo4J offers two interfaces - a lightweight declarative query language called *Cypher*, and a Neo4J offers two interfaces - a lightweight declarative query language called *Cypher*, and a Java API more appropriate for complex graph algorithms. PostgreSQL only offers one interface through SQL. Finally, Grapht takes the same approach as Neo4J in offering both a programmatic API and a declarative query language based on SQL called *gSQL*.

When considering the cost of transitioning to a new database system, the time taken to adopt to use the new system should also be examined. This cost of transition is one area in which Grapht particularly shines, since no changes to the underlying database need be made, and since the *gSQL* query language is a relatively small extension to standard SQL. It is difficult to comment on the ease of use of the different query languages without a full usability study. However I would expect to find that neither the Grapht nor Neo4J APIs would score highly in this kind of study. Instead, I would expect that *Cypher* perform well for novices, since its pattern-matching queries are expressed in a fairly intuitive visual language. On the other hand, SQL is much more likely to have been encountered before by experienced developers, which would presumably lead to a higher score. *gSQL*, as a relatively small extension on top of SQL, would likely be found easier to use by experienced developers than by novices.

It is clearly desirable from an ease-of-learning perspective to only need a single interface. One might first explore why it is that PostgreSQL can afford to take this approach, while the more graph-oriented approaches must provide a lower-level API alongside a declarative query language. The answer lies in considering the relative search spaces of the two problems. Various optimisations and heuristics are possible to reduce the cost of traversing a tabular relation, but in the worst case it is always possible to retrieve a result in linear time in the number of rows by enumerating each row. The same is not true when considering paths through a graph. The number of possible paths grows exponentially with each additional hop in the path, meaning that for most problems it is not practical to simply enumerate all possible paths when searching for a particular one. Instead, an intelligent approach is required to direct traversal through the graph.

Cypher, having been in development for some time, has the ability to exploit certain statistics and heuristics in order to slightly cut down the search space, but does not avoid the problem in any real way. To find the shortest path between two points using Cypher, for example, one must enumerate all possible paths between the points, and rank them by increasing length. To avoid this, a large number of built-in functions are provided (such as `shortestpath`, which can be used to avoid this situation) which can internally specify a particular traversal strategy. However these built-in functions do not truly address the shortcoming of the language, which cannot be used to express custom algorithms for which a built-in has not been provided. When using Grapht's or Neo4j's traversal API directly, the user implements her own traversal strategy, removing the guesswork required for a declarative interface. gSQL sacrifices some of the declarative nature of the query language by allowing users to specify a traversal strategy with a `TRAVERSE BY` clause. This allows solutions to be found for a query more quickly, though at the cost of slightly increased development complexity.

A second factor is the ability to create and manipulate data structures. Almost all algorithms of interest manipulate as part of their operation some central data structure. It is often difficult to define these structures using a

declarative query language, which typically tries to abstract away implementation details such as the choice of data structure for an operation. With some creativity, it is often possible to manage a workaround in SQL, by using temporary tables to encode the data structures required. A priority queue can be implemented in this way (as required for a full-SQL implementation of A\*), though it will rarely perform as well as a solution using binary heaps. This is not possible at all in Cypher, which can perform little more than pattern-matching on paths through the graph. The path searches possible through gSQL are directed using the **TRAVERSE BY** clause, which internally uses a priority queue to select upcoming expansions. Much as it is possible to implement other data structures using a SQL table, it is also possible to make creative use of Grapht’s priority queue to implement some other structures. For example, a “last-in, first-out” stack can be created by using insertion time as priority. This is not quite so flexible as SQL, but again allows slightly more complicated algorithms to be implemented entirely in gSQL than is possible in a language like Cypher.

Although some queries are totally inexpressible in a particular query language, the ease of development for expressible queries should also be considered. This is where Neo4J and Grapht’s traversal APIs fall down. Although any possible query can be expressed using the full power of any JVM language, needing to prepare and compile code in order to express a simple query can be a huge complication, especially when much of the code will be concerned with routine tasks like establishing a database connection or deserialising data. Although Java is a very well-known programming language, using APIs in this way also requires that users be familiar with Java as well as the database API itself. For newcomers, this means overcoming two learning barriers rather than just one. In some senses, the dual interfaces offered by Neo4J and Grapht alleviate this problem, by providing a lightweight declarative query language for simple queries, while providing a more powerful API which can be used for more complex needs. However introducing a second interface is also likely to create confusion about which one to use in different contexts, and at the very least increases the amount

of learning required to effectively use the system to its full potential.

# Chapter 5

## Conclusion

TODO





# Bibliography

- [1] E. F. Codd. “A Relational Model of Data for Large Shared Data Banks”. In: *Commun. ACM* 13.6 (June 1970), pp. 377–387. ISSN: 0001-0782. DOI: 10.1145/362384.362685. URL: <http://doi.acm.org/10.1145/362384.362685>.
- [2] Peter Pin-Shan Chen. “The Entity-relationship Model—Toward a Unified View of Data”. In: *ACM Trans. Database Syst.* 1.1 (Mar. 1976), pp. 9–36. ISSN: 0362-5915. DOI: 10.1145/320434.320440. URL: <http://doi.acm.org/10.1145/320434.320440>.
- [3] Jim Gray et al. “The transaction concept: Virtues and limitations”. In: *VLDB*. Vol. 81. 1981, pp. 144–154.
- [4] Serge Abiteboul and Nicole Bidoit. “Non First Normal Form Relations to Represent Hierarchically Organized Data”. In: *Proceedings of the 3rd ACM SIGACT-SIGMOD Symposium on Principles of Database Systems*. PODS ’84. Waterloo, Ontario, Canada: ACM, 1984, pp. 191–200. ISBN: 0-89791-128-8. DOI: 10.1145/588011.588038. URL: <http://doi.acm.org/10.1145/588011.588038>.
- [5] H. S. Kunii. “DBMS with Graph Data Model for Knowledge Handling”. In: *Proceedings of the 1987 Fall Joint Computer Conference on Exploring Technology: Today and Tomorrow*. ACM ’87. Dallas, Texas, USA: IEEE Computer Society Press, 1987, pp. 138–142. ISBN: 0-8186-0811-0. URL: <http://dl.acm.org/citation.cfm?id=42040.42071>.

- [6] Joachim Biskup, Uwe Rasch, and Holger Stiefeling. “An extension of SQL for querying graph relations”. In: *Computer Languages* 15.2 (1990), pp. 65–82.
- [7] W. Kim. “Object-Oriented Databases: Definition and Research Directions”. In: *IEEE Trans. on Knowl. and Data Eng.* 2.3 (Sept. 1990), pp. 327–341. ISSN: 1041-4347. DOI: 10.1109/69.60796. URL: <http://dx.doi.org/10.1109/69.60796>.
- [8] Grzegorz Malewicz et al. “Pregel: a system for large-scale graph processing”. In: *Proceedings of the 2010 ACM SIGMOD International Conference on Management of data*. ACM. 2010, pp. 135–146.
- [9] Amit Singhal. “Introducing the knowledge graph: things, not strings”. In: *Official Google Blog, May* (2012).
- [10] Eiko Yoneki and Amitabha Roy. *A unified graph query layer for multiple databases*. Tech. rep. UCAM-CL-TR-820. Department of Computer Science, University of Cambridge, 2012.
- [11] Julian Shun and Guy E Blelloch. “Ligra: a lightweight graph processing framework for shared memory”. In: *ACM SIGPLAN Notices*. Vol. 48. 8. ACM. 2013, pp. 135–146.
- [12] Robert Campbell McColl et al. “A performance evaluation of open source graph databases”. In: *Proceedings of the first workshop on Parallel programming for analytics applications*. ACM. 2014, pp. 11–18.
- [13] Wen Sun et al. “SQLGraph: An Efficient Relational-Based Property Graph Store”. In: *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data*. SIGMOD ’15. Melbourne, Victoria, Australia: ACM, 2015, pp. 1887–1901. ISBN: 978-1-4503-2758-9. DOI: 10.1145/2723372.2723732. URL: <http://doi.acm.org/10.1145/2723372.2723732>.
- [14] R Pointer. “Introducing FlockDB, Twitter, Inc.” In: *Available: <https://blog.twitter.com/flockdb>* ().